

Package CIpostSelect

Boubacar DIALLO

November 10, 2025

Abstract

Traditional or classic methods and the PoSI method for post-selection inference have limitations, particularly inadequate confidence intervals in some situations, highlighting the need for a more precise approach. Our method is based on repeated data splits, along with two voting systems.

The results obtained from this methodology in simulations show empirical validation of the confidence intervals. Additionally, in terms of performance for identifying the true model in a variable selection procedure, it outperforms Lasso and BIC when used alone.

1 Motivations

Today, with the availability of massive datasets, we often encounter many predictor variables, not all of which are relevant. This leads us to perform variable selection to identify a subset of variables that best explains our target. However, making inferences or calculating confidence intervals after such variable selection is challenging, especially due to the bias introduced by the selection process.

Studies, such as the one by Zhang et al. (2022), have shown through simulations that classical methods ($\text{Algo}_{\text{selection}} + \text{OLS}$) tend to produce confidence intervals that are too narrow, as illustrated in the table below:

D. Zhang et al.

TABLE 1
Empirical coverage probabilities of 95% naïve confidence intervals for the true non-zero β_j^0 's.

Model Selector + Estimation	Empirical coverage probabilities				
	β_1^0	β_3^0	β_4^0	β_5^0	β_7^0
LASSO + OLS	.691	.486	.695	.736	.560
Elastic net + OLS	.749	.493	.815	.867	.619
SCAD + OLS	.362	.736	.682	.727	.723

In this study, starting with a larger set of variables M such that $|M| > 5$ and $Mo = 5$ where $Mo \subset M$, they repeated simulations to calculate the confidence intervals 95% as follows:

$$\text{CI}_{\hat{\beta}_j^{\hat{M}}} = \left[\hat{\beta}_j^{\hat{M}} - t(n - |\hat{M}|, 1 - \frac{\alpha}{2}) \times \text{sd}(\hat{\beta}_j^{\hat{M}}), \hat{\beta}_j^{\hat{M}} + t(n - |\hat{M}|, 1 - \frac{\alpha}{2}) \times \text{sd}(\hat{\beta}_j^{\hat{M}}) \right] \quad (1)$$

After each simulation, they checked if the true coefficients ($\beta_1^0, \beta_3^0, \beta_4^0, \beta_5^0$, and β_7^0) were within their respective confidence intervals.

For a good post-selection confidence interval method, we expect probabilities around 0.95. However, in some simulations, we see that classical methods give confidence intervals that are too narrow.

Furthermore, we explored a popular post-selection inference method called **PoSI methods**. This method is described as providing valid confidence intervals for any sub-model in a larger model.

The idea is to study all possible submodels of our initial model M , to get an arbitrary K that replaces the Student distribution quantile in expression (1). This K is chosen so that:

$$\mathbf{P}(\max_{j \in \hat{M}} |t_j^{\hat{M}}| \leq K_{\text{PoSI}}) \geq 1 - \alpha \quad (2)$$

where $t_j^{\hat{M}}$ is the test statistic associated with the variable j .

The following graphs show that PoSI methods (2) tend to give much wider confidence intervals, with an empirical coverage probability equal to 1.

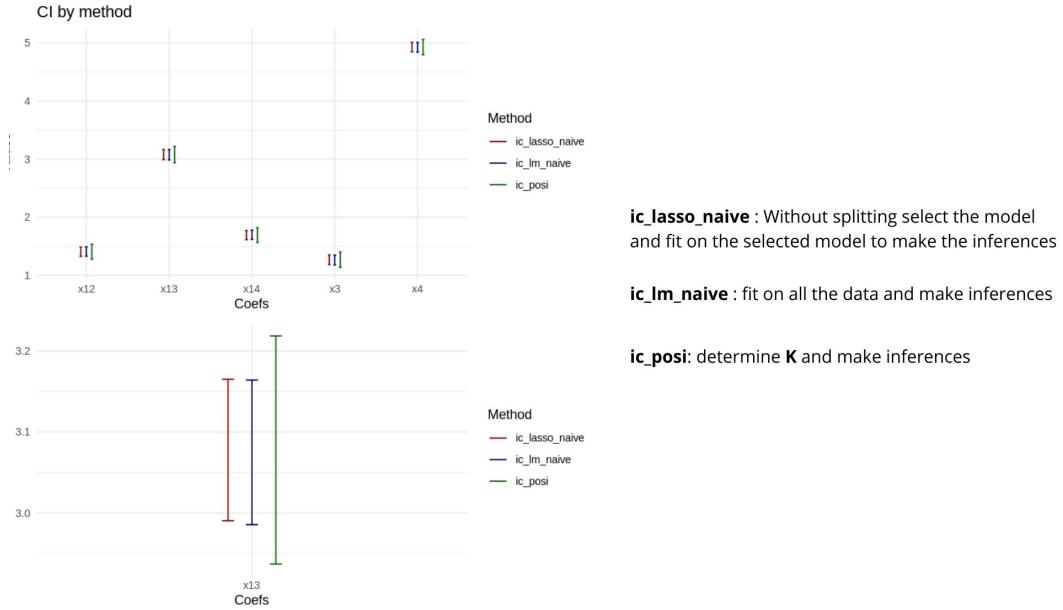


Figure 1: Comparison of Confidence Interval Widths with Three Methods:

We conducted a simulation with $p = 15$ and $p_0 = 5$. Using these simulated data, we calculated confidence intervals with the three methods described on the right of the figure.

We observe that the confidence intervals produced by PoSI are much wider.

$p = 15, p_0 = 5$

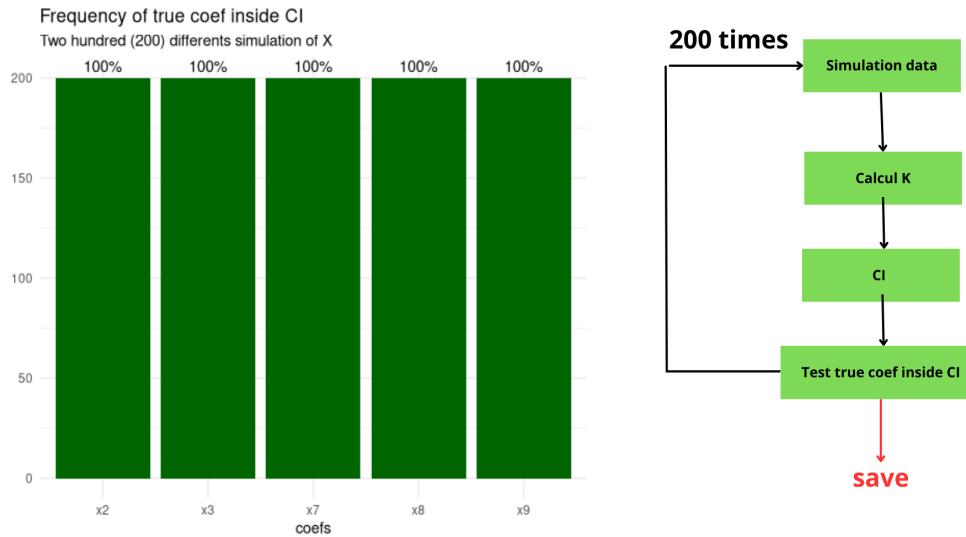


Figure 2: Empirical validation of PoSI confidence interval:

We observe that PoSI is too conservative, as we chose K at 95% in 200 simulations, and the true coefficients are always within the constructed confidence intervals.

Furthermore, since PoSI examines all sub-models, it should be noted that when the number of explanatory variables exceeds 15 ($p > 15$), it takes a considerable amount of time to obtain K .

Following these observations, we decided to develop a method that would fall between the two, producing appropriately sized confidence intervals, neither too narrow nor too wide.

2 Methodology

Our method relies on repeated data splitting with two types of voting, which we will present below.

2.1 Data splitting

- First step : data splitting

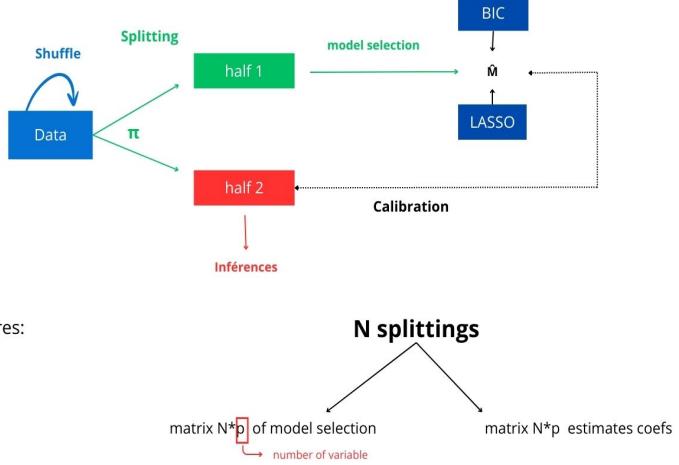


Figure 3: **Data splitting strategy**

Given our data, for a certain number N of repetitions, we shuffle the data and split it into two parts with a probability π . On one part, we perform model selection $\hat{M}^{(k)}$, and on the other part, we calibrate the model to obtain coefficient estimates $\hat{\beta}^{(k)}$ with $k \in [1, \dots, N]$. As a result, we end up with an $N \times p$ matrix for model selections and an $N \times p$ matrix for the corresponding coefficient estimates for p explanatory variables.

The next question is how to choose a final model from this matrix of models. Once the model is selected, the following challenge is determining how to compute confidence intervals and obtain coefficient estimates corresponding to the chosen model.

2.2 Voting system

We will explore two types of voting: one on the coefficients and one on the models. Each of them has a specific way of obtaining the estimates.

2.2.1 Vote on models

For the voting on models, we proceed as follows:

$$\hat{M}_{\text{vote}} = \hat{M}^{(j)} \quad \text{s.t. } \forall i \neq j : \#\hat{M}^{(j)} > \#\hat{M}^{(i)}$$

We choose the model recorded during the N splittings that has been selected the most frequently. Next, we determine

$$A = \{k \in \{1, \dots, N\} : \hat{M}^{(k)} = \hat{M}_{\text{vote}}\}$$

Then, we compute the coefficient estimates as follows:

$$\hat{\beta}_j^{\hat{M}_{\text{vote}}} = \frac{1}{|A|} \sum_{k \in A} \hat{\beta}_j^{\hat{M}^{(k)}} \quad \forall j \in [1, \dots, p]$$

We will then search for the rows corresponding to the voted model in the coefficient estimates and average them across the columns to obtain the coefficient estimates, as illustrated in the following graphic :

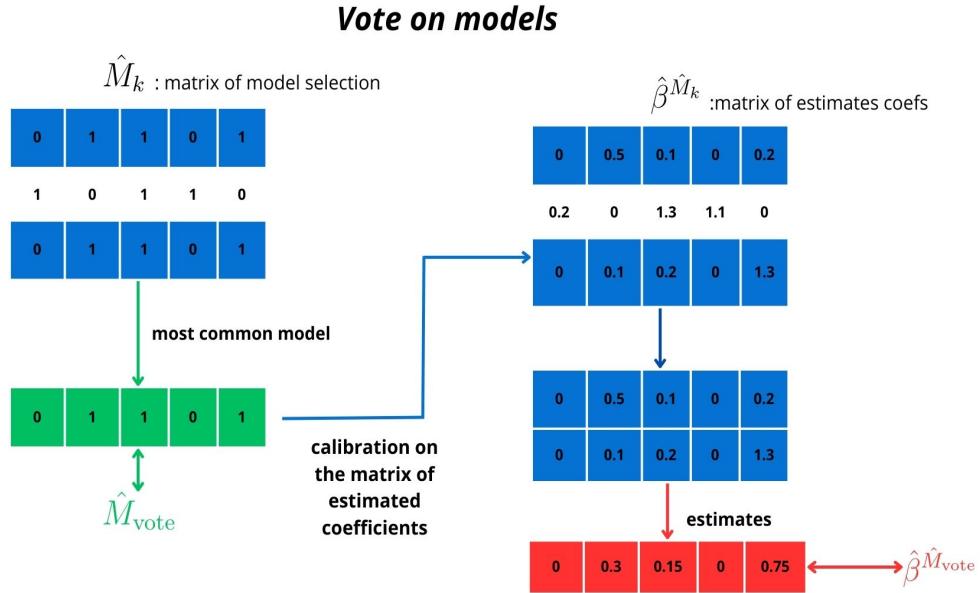


Figure 4: **Example of voting on models with $N = 3$ splittings and $p = 5$ explanatory variables**

2.2.2 Vote on the coefficients

For the voting on coefficients, we will proceed as follows:

We will start by selecting the model by checking for each variable whether it has been selected a certain percentage of the time, for example, 50%.

$$\hat{M}_{\text{vote}} = \left\{ \begin{array}{l} \frac{1}{N} \sum_{i=1}^N \hat{M}_{ij} > s \\ j \in [1, \dots, p] \end{array} \right\}$$

where s is the threshold we set to decide whether or not to select the coefficient. Next, we will focus exclusively on the indices corresponding to the selected variables, which we will call index J :

$$J = \{j \in P = \{1, \dots, p\} \quad / \quad \hat{M}_{\text{vote}}^P = 1\}$$

We will obtain the estimates by averaging the rows where the retained variables were selected:

$$(\hat{M}_k)_J = \begin{cases} 0_N & \text{if } j \notin J \\ \hat{M}_k^{(j)} & \text{if } j \in J \end{cases} \quad \forall j = 1, \dots, p$$

For all $j \in P$:

$$a_j = \{k : (\hat{M}_k)_J^{(j)} = 1\}$$

Then, we compute:

$$\hat{\beta}_j = \frac{1}{|a_j|} \sum_{a_j} \hat{\beta}_j^{\hat{M}_k}$$

All of this procedure is summarized in the example below:

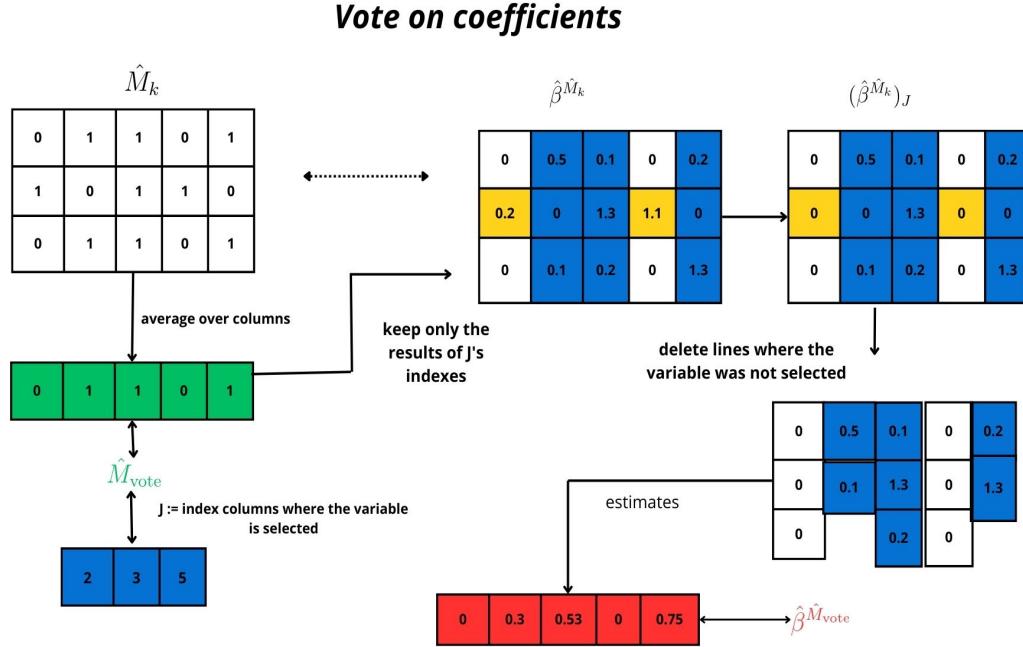


Figure 5: Example of voting on the coefficients with $N = 3$ splittings and $p = 5$ explanatory variables

2.3 Empirical method

I spoke only about coefficient estimates and not about calculating confidence intervals. It is important to note that at the end of the N splittings, with N large, we will have, as seen in the last motifs of our examples before averaging to obtain the coefficients, each column representing a vector of coefficient estimates for a given variable.

Thus, for the confidence intervals, we will then take the empirical quantile for each coefficient vector, as illustrated in the graphic below:

Empirical method

- No theoretical knowledge
- Resampling

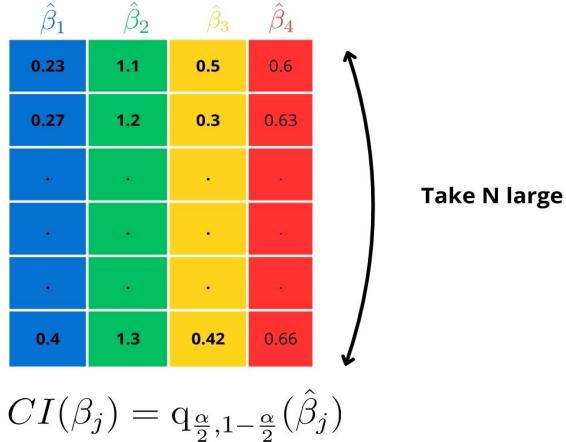


Figure 6: Empirical method for confidence intervalle

3 CIpostSelect

After this work on the methodology, we developed the **CIpostSelect** package, also known as "*Confidence Interval Post Selection of variables*", which is available on CRAN.

This package features two main functions:

- `lmps` provides the model selection matrices from the various splittings of our data, as well as the matrix of coefficient estimates for the selected models during the different splittings.
- `CIps` takes an `lmps` object as a parameter, along with other parameters that specify the type of voting, for example, and the confidence level of our confidence intervals.

For more details, please refer to the documentation: `help("CIpostSelect")`

3.1 Example using a dataset in R

We will apply our package to the **BostonHousing** dataset.

The idea here is not to interpret the results but to show you how to apply the package. We will revisit the performance of the package a little later.

BostonHousing:

The BostonHousing dataset, from the `mlbench` package in R, contains 506 observations and 14 variables that describe characteristics that influence house prices in Boston. Variables include socioeconomic and environmental factors, such as the crime rate per capita, the number of rooms, the distance to employment centers, and the property tax rate. The target variable, `medv`, represents the median value of owner-occupied homes in thousands of dollars.

```
library(CIpostSelect)
library(mlbench) # BostonHousing dataset
library(doParallel) # for parallelization
# Load the BostonHousing dataset
data("BostonHousing")
# Fit model --> lmps object
model = lmps(medv ~ ., data = BostonHousing, method = "Lasso", N =
500, p_split = 0.5, cores = detectCores() - 2)
```

We recorded in the `model` the results of model selection as well as coefficient estimates from the $N = 500$ splittings of our BostonHousing data set with a splitting probability of $\pi = 0.5$ and a post-selection Lasso.

```
> summary(model)
Execution time in seconds:
[1] 10.5

Number of Splits
[1] 500

Frequency of variable selection :
  crim zn indus chas nox rm age dis rad tax ptratio b lstat
Nb   193   57     43   303   124   500    15   183    14   165   500   379   500

The number of times the most frequent model was selected :
[1] 65
```

The results of calling the `summary` function are very important, as they will inform us later on which voting strategy is best suited to calculate confidence intervals.

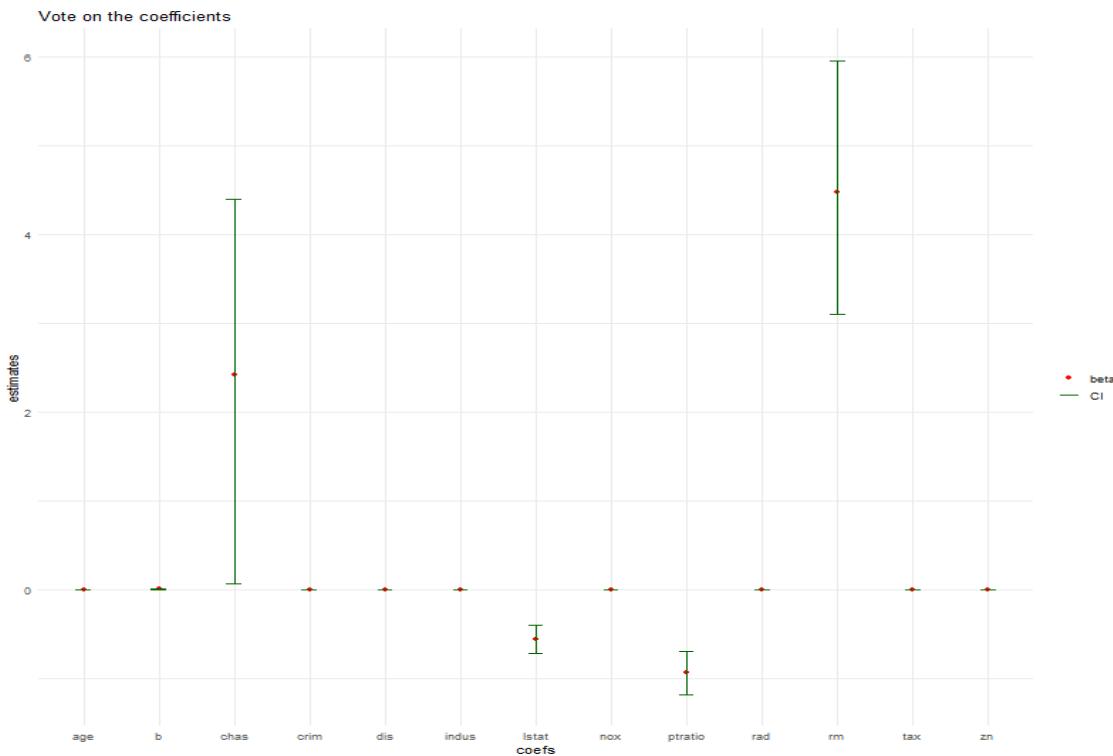
In this example, we notice that the number of times we find the most frequent model is quite low compared to the number of divisions we performed. This is why we will opt for a vote based on the coefficients.

```
cips = CIps(model, vote = "coef", alpha = 0.05, s.vote_coef = 0.5)
```

Thus, we decide to perform a vote on the coefficients for confidence intervals at 95% ($\alpha = 5\%$) with a selection threshold of $s = 0.5$.

```
> print(cips)
Variable selection :
  crim   zn  indus  chas   nox    rm   age   dis    rad   tax  ptratio    b   lstat
selected? FALSE FALSE FALSE TRUE FALSE TRUE FALSE FALSE FALSE TRUE TRUE TRUE TRUE
Estimates :
      chas        rm      ptratio        b      lstat
lower  0.0701    3.0192   -1.1788    0.0045   -0.7110
upper  4.2346    5.9552   -0.6922    0.0150   -0.3544
beta   2.4262    4.4975   -0.9247    0.0098   -0.5604

> plot(cips)
```



As I mentioned earlier, we will not focus on interpretations here; this was simply an example to explain how to use the `CIpostSelect` package to calculate post-selection confidence intervals for variable selection. We will now discuss the performance of the package, particularly in comparison to classical methods.

4 Performance – Confidence Interval

In this section, we will study a simulation very similar to that of Zhang et al. (2022), which showed that the confidence intervals provided after a Lasso post-selection on the entire dataset tended to be very narrow and did not achieve their empirical coverage probability.

We will use this type of simulation to first compare the **width of the confidence intervals** of our different **methods**(classics versus `CIpostSelect`), and then examine their **empirical coverage**.

Notation:

- "**CIpostSelect**" := will denote a BIC post-selection strategy

- "*CIpostSelect2*" := will denote a Lasso post-selection strategy

Simulation:

- $n = 60, p = 18, po = 5$

We have 60 observations of 18 explanatory variables, of which 5 are relevant (true non-zero coefficients).

- $\beta_{j \in [po]} = 1.5$

The relevant coefficients are all equal.

- $X \sim \mathcal{N}(0, 1)$ and $\epsilon \sim \mathcal{N}(0, \sigma^2)$ $\sigma^2 = 2$

Methods:

- Classical methods (Lasso or BIC + OLS)

Selection on the entire dataset and inference on the entire dataset while restricting to the selected subset.

- CIpostSelect

4.1 Width of Confidence Intervals

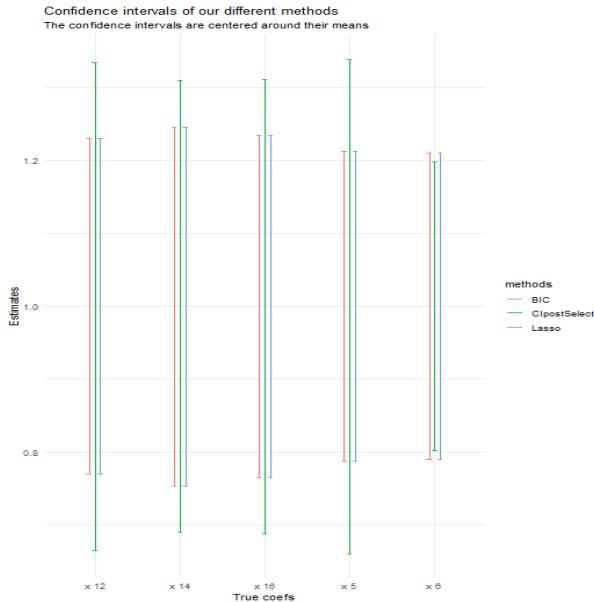
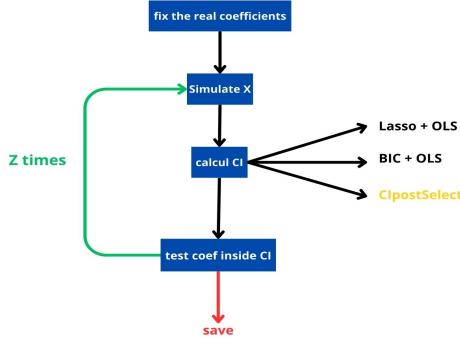


Figure 7: **Width of Confidence Intervals of Our Different Methods:**

We can notice that the confidence intervals provided by our method, CIpostSelect, shown in green, are much wider than those of classical methods.

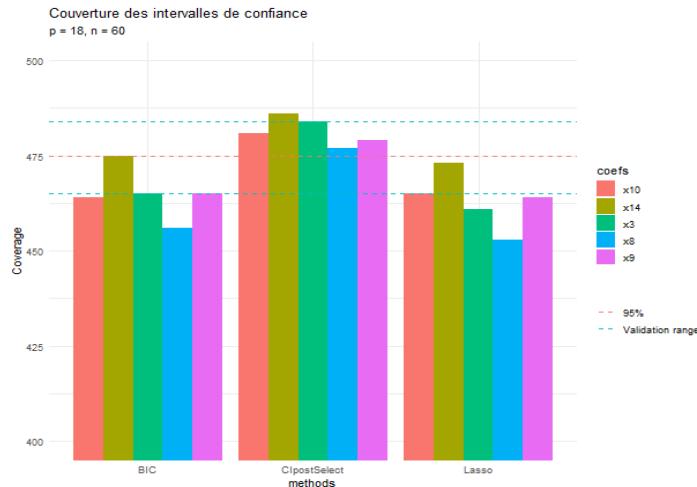
4.2 Empirical coverage

Now we move on to the empirical coverage probability of the confidence intervals. We will repeat the calculation of confidence intervals a certain number of times, Z times, and at each repetition, we will test whether the true coefficients $\beta_{[j \in po]}$ are within the confidence intervals of the different methods.

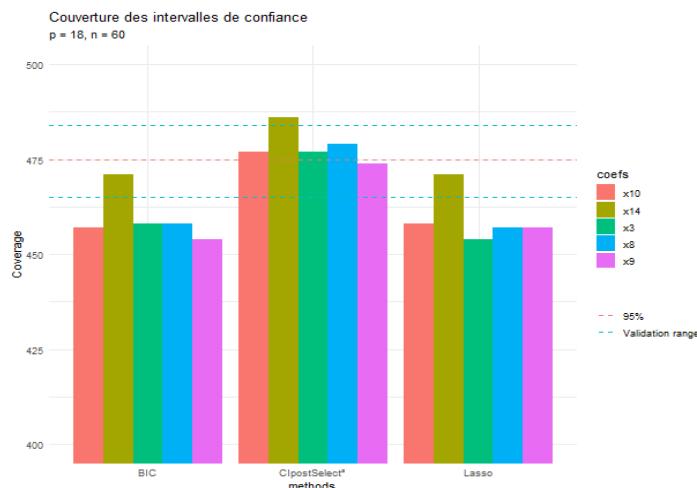


We will also define a validation range for the 95% confidence intervals, which is given by:

$$[q_{\frac{\alpha}{2}}(\mathcal{B}(Z, 0.95)), \quad q_{1-\frac{\alpha}{2}}(\mathcal{B}(Z, 0.95))]$$



(a) BIC post-selection strategy for CIpostSelect



(b) Lasso post-selection strategy for CIpostSelect

Figure 8: Empirical coverage of our methods:

We can notice that our method for calculating post-selection confidence intervals from our package is empirically validated, in stark contrast to classical methods.

5 Performance – Selection of the true model

This is not the primary goal of this project development. However, we observed that throughout the simulations and tests on the methodology, we were able to accurately identify the true model M_0 . This prompted us to test the model selection performance of our package `CIpostSelect` compared to classical methods.

The strategy will be the same as that adopted for empirical coverage. The only difference here is that we will stop at model selection, and at this stage, we will test in each simulation whether the model we voted for using our various methods is the same as the fixed model M_0 .

Independent Simulation

- $n = 300, p = 50, po = 15$
- $\beta_{[j \in po]} \sim \mathcal{U}[1, 5]$
- $X \sim \mathcal{N}(0_p, I_p)$

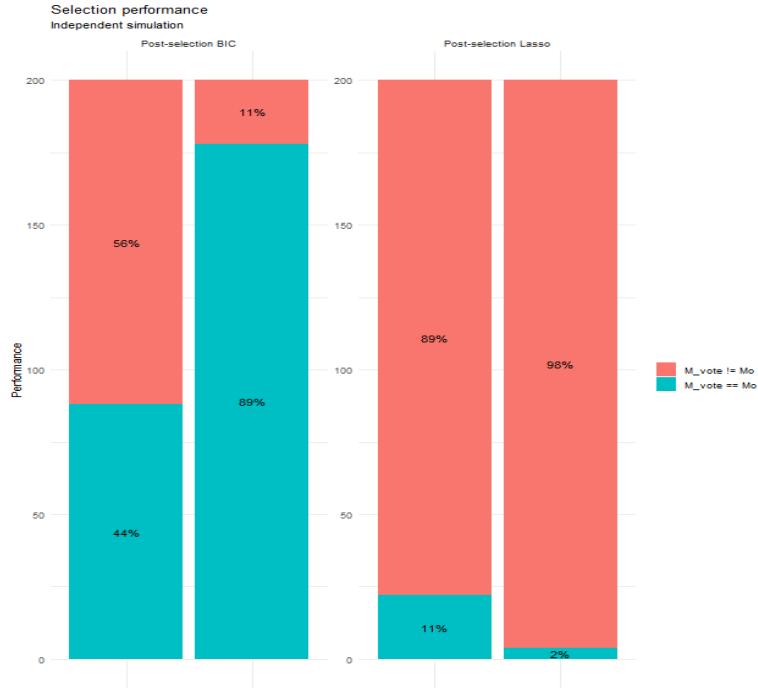
Dependent Simulation

We have the same parameters as in the independent simulation, but this time we introduce dependence among the variables as follows:

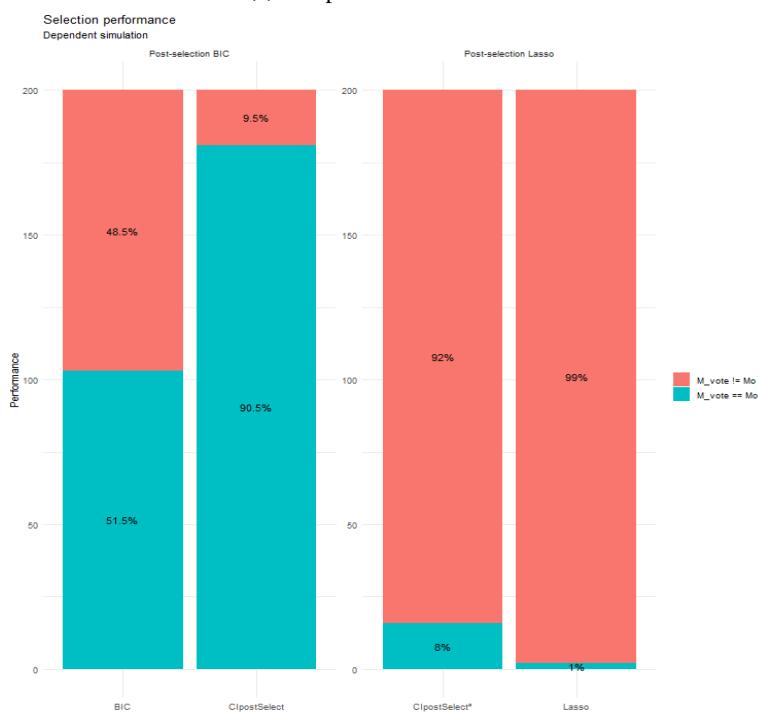
- $X \sim \mathcal{N}(0_p, \Sigma)$
- $\Sigma = \begin{bmatrix} I_{20} & & \\ & I_{20} + \alpha A_{20} & \\ & & I_{10} + \gamma A_{10} \end{bmatrix} \in \mathbb{R}^{p \times p}$ with $\alpha = 0.1$ and $\gamma = 0.5$

Thus, we have dependent data, and we ensured that among the relevant variables, there are independent ones, weakly dependent ones ($\alpha = 0.1$), and strongly dependent ones ($\gamma = 0.5$).

To test the selection performance for each of the two types of simulations, we will conduct $Z = 200$ simulations of our data, and the frequency of retrieving the true model is given by the results below:



(a) Independent Simulation



(b) Dependent Simulation

Figure 9: Selection Performance of Our Methods on Two Types of Simulations:

We can observe that our package provides a significant improvement in selecting the true model compared to classical methods.

For example, with dependent data using a BIC post-selection strategy, our package CIpostSelect retrieves the true model 90% of the time, while BIC alone only succeeds 51% of the time.

6 CIpostSelect-advanced: force a variable

In a project involving variable selection to address a specific issue, such as the *impact of antibiotics on children's growth*, it may happen that the variable of interest (here, the use of antibiotics) is not selected by standard variable selection methods. This might suggest that, from the perspective of explaining the target (children's growth), the use of antibiotics is not as significant as other variables, like the child's age, birth weight, etc.

The question then arises: **how do we address the issue if the variable of interest is not selected?** One solution is to force the selection of this variable. In our case, this means forcing the constant selection of the *use of antibiotics* variable during the N splits.

What remains to be determined is whether the resulting inferences are valid. To do this, we will force, in two types of simulations (independent and dependent, as before), the selection of variable "X1" under the following conditions:

- $\beta_1 = 0.2$
- $\beta_{[j \neq 1 \in po]} \sim \mathcal{U}[1, 5]$

Thus, this variable "X1" is less influential than the others, making it more difficult to select under normal conditions. We will then observe the resulting outcomes.

To force a variable, you need to specify the name of the variable you want to force when calling the `lmps` function, as shown below:

```
model = lmps(y ~ ., data = data, method = "BIC", N = 1000, p_split =  
0.5, cores = detectCores() - 2, forced_var = "X1")
```

Independent Simulation:

- **Number of variables (p):** 18, **Number of relevant variables (po):** 5
- $\beta_1 = 0.2$ and $\beta_{[2,po]} \sim \mathcal{U}[1, 5]$
Simulate the 4 relevant variables among the 5 uniformly from [1, 5], with the coefficient of the forced variable $X1$ set to 0.2.
- $X \sim \mathcal{N}(0, 1)$ and $\epsilon \sim \mathcal{N}(0, 2)$
- Sample sizes n in [60, 100, 150, 200]

For each of these simulations, we will run **CIpostSelect** and **CIpostSelect²** with each of their two votes to compare the estimates of the forced variable.

Dependent Simulation: The same approach will be used, but we will add dependence between the variables as described in **section 4**.

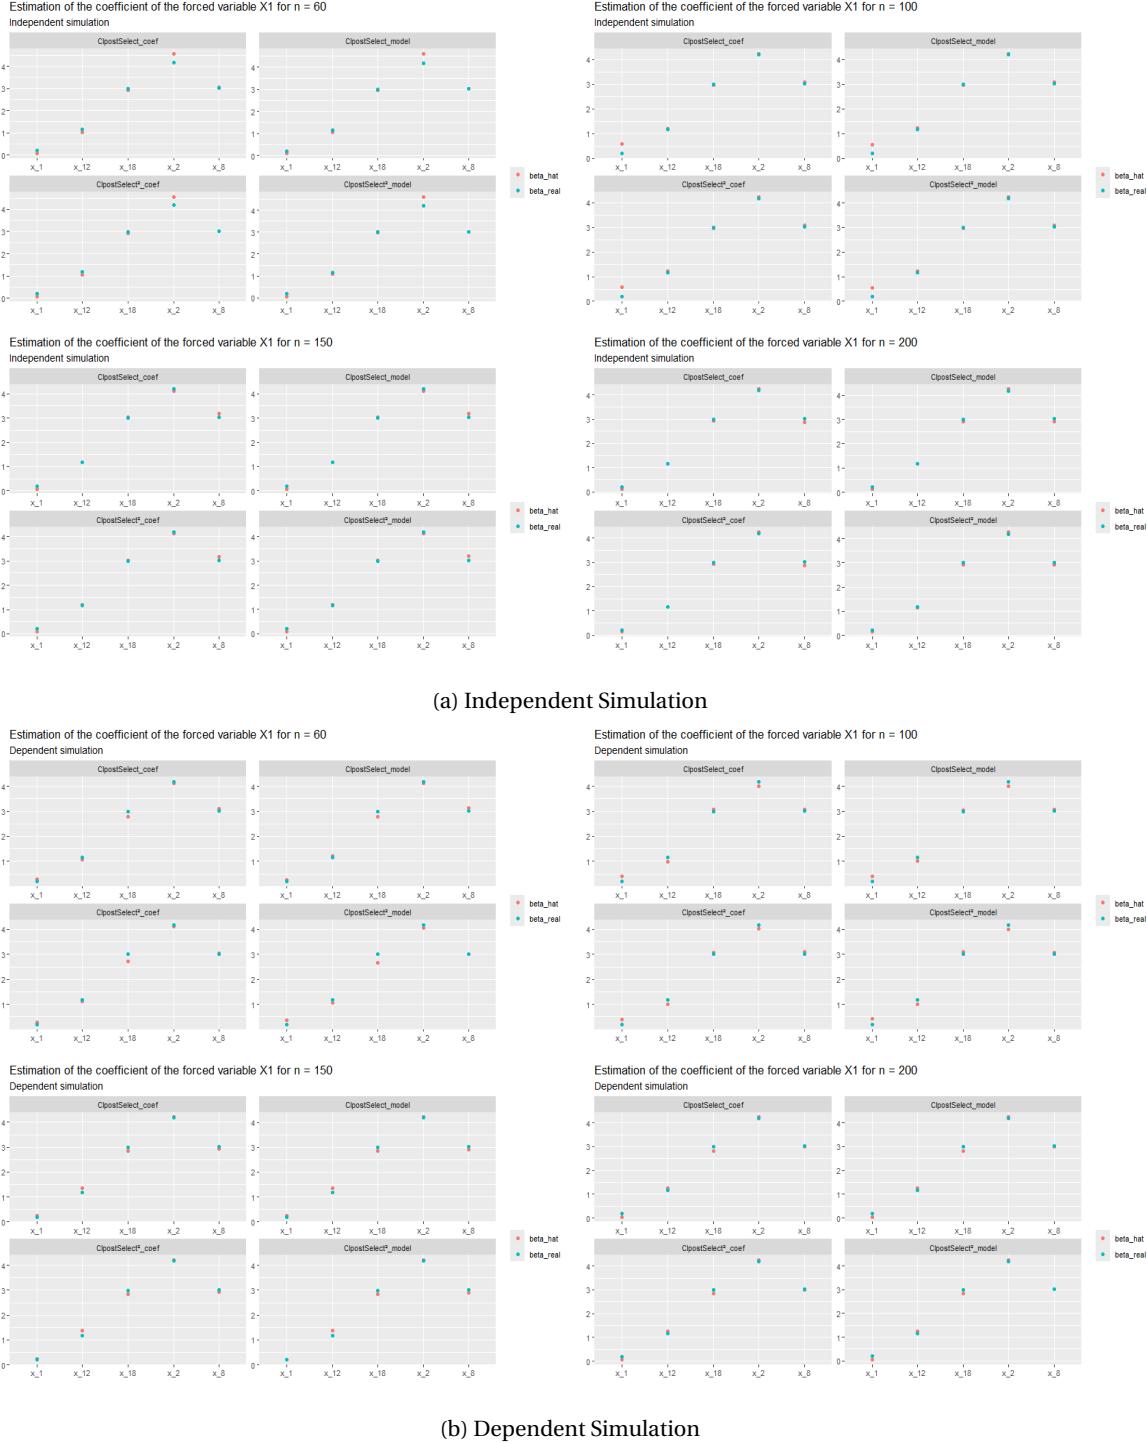


Figure 10: Estimation of the coefficient of the forced variable X_1

We can observe that the estimated coefficients for the variable X_1 across different methods and simulations are close to the true value of the actual coefficient. This gives us confidence in using this option when the variable of interest is not selected and we want to interpret it.